

Singularity containers and SLING a tutorial

Blaž Škrlj

Jožef Stefan Institute

Purpose of this tutorial

1. Get to know one of the common pitfalls of research code
2. How to mitigate this issue via dockerization
3. What are Singularity Containers?
4. **Hands-on:** Local builds
5. Cluster execution
6. **Hands-on:** Scaling up with SLING

Part 1

Code replicability and Singularity

The four levels of repositories one finds these days

1. Messy, undocumented code with *no dependency specifications*
2. Code with dependency specifications, *not versioned*
3. Versioned code, dependencies
4. **Code with a replicable environment**

(2) What is Singularity?

1. It's an environment for **dockerization**. Specifically focused on **scientific** and **HPC** environments.

In other words:

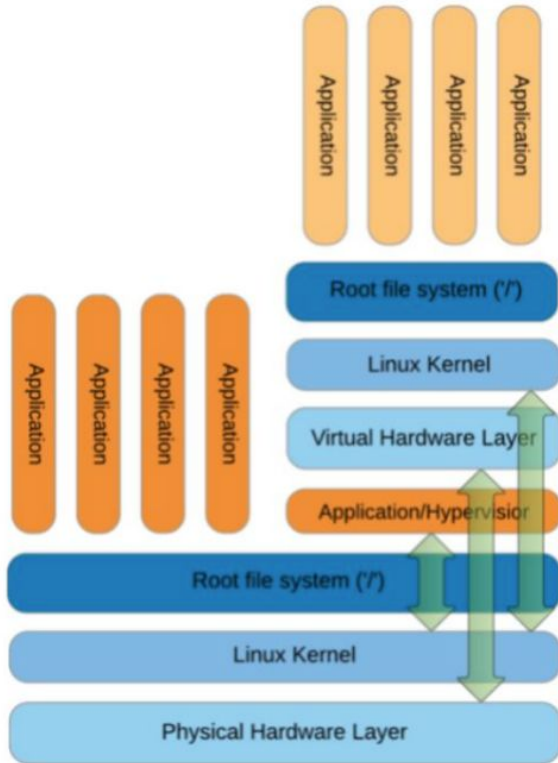
It's the engineering swiss knife suitable for the **fast-paced** (replicable) science.

(1) What is Singularity?

1. Assuming you wish to share some code with a colleague
2. What's the minimal effort she/he needs to undergo to run the code?




Virtual Machine



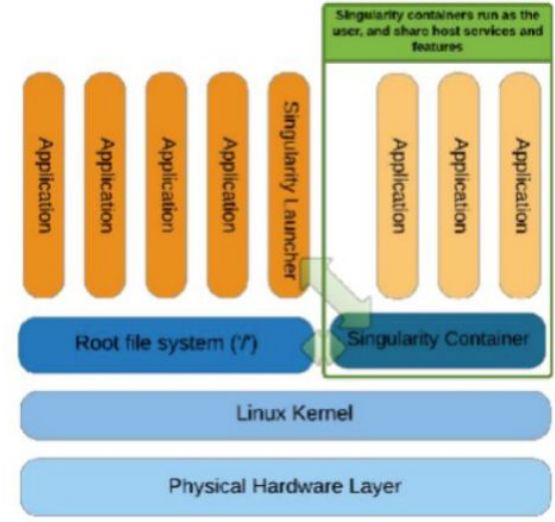
Docker

- Independent namespaces
- Root supported contexts



Singularity

- Host services and devices shared
- User access



Summary

When **to use** it

Testing

- Unstable releases
- Proof of concept executions
- Punctual data preproces

Reproducibility

- Complex software stacks

When **NOT to** use it

Production

- Architecture aware software
GROMACS, NAMD, CPMD...
- Frequent HPC applications

Not available soft?

- Request it to the support team

Detailed info at: <https://sylabs.io/docs/>

Extremely useful for

Deployment on **multiple machines** (HPC).

Quick **replicable** demos and handling dependency problems.

“Oh, this machine is not doing anything, let’s just run this experiment” - type of adventures

Tutorial time - part 1

1. Installation
2. Simple generic replicability case
3. Building containers
4. Using GPUs

Available at: **<https://github.com/SkBlaz/singularity-container-tutorial#enter-singularity>**

Part 2

Singularity and SLING

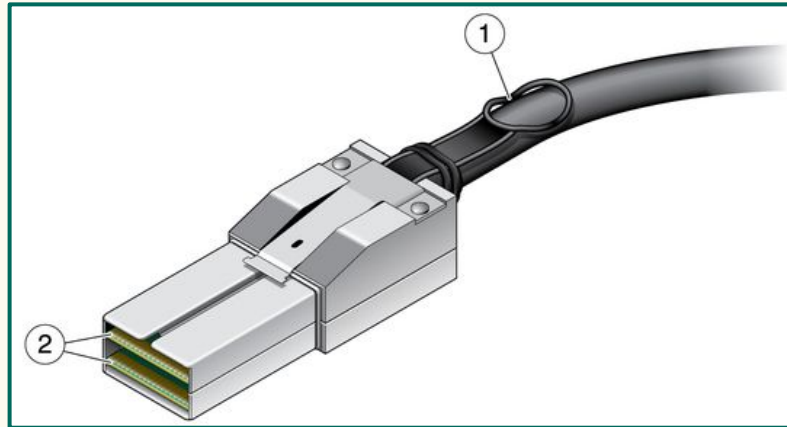
SLING?

Site	CPU	Load (processes: Grid+local)
ARNES	2048	1845+8
CIPKeBiP	984	144+0
Maister	9872	200+8794
NSC	2880	1424+540
SiGNET	8096	3669+3804
SiGNET	8096	3804+3669
Trdina	536	100+16
UNG	240	232+0
8 sites	32752	11418 + 16831

Learn more at: <http://www.sling.si/sling/en/eurocc/>

Why Singularity in this environment?

1. Special support (InfiniBand)
2. No need for customized environments (SysAdmin support for each task)
3. Fast



Source: <https://docs.oracle.com/cd/E19654-01/835-0786-02/figures/127163.jpg>

Key idea

1. Build the Singularity image locally, copy it to dCache
2. Each job sees the image and executes it (singularity exec [jobStuff])
3. The dCache version of the image can be updated by the user

If it works locally, it will **most likely** work on the grid too.

Tutorial time - part 2

1. Let's build an image
2. Test it locally
3. Create XRS� job specification
4. Execute, retrieve and inspect the result

Available at:

<https://github.com/SkBlaz/singularity-container-tutorial/blob/master/gpu.md>

<https://github.com/SkBlaz/singularity-container-tutorial/blob/master/optimization.md>